

aucmat: matrix-first ROC analysis for high-dimensional biomarker screening in R

Lucas VHH Tran^{1,*}

¹CK-CARE, Davos, Switzerland

*lucas.tran@ck-care.ch

Abstract

Summary: The area under the receiver operating characteristic curve (AUC) is the standard metric for evaluating continuous biomarkers against binary outcomes. When screening thousands of molecular biomarkers, researchers need matrix-wide AUC estimation with proper inference, multiplicity adjustment, and stability assessment. We present **aucmat**, an R package providing a matrix-first architecture: `aucmat(X, y)` computes direction-preserving AUCs with DeLong or stratified-bootstrap confidence intervals for every column of a numeric matrix in a single call. The package supports Benjamini-Hochberg multiplicity adjustment, three hypothesis-testing frameworks (superiority, non-inferiority, equivalence via TOST), an omnibus Wald test for equality of multiple correlated AUCs via the joint DeLong covariance matrix, and bootstrap rank-stability analysis. A two-stage Hurdle-AUC model handles zero-inflated data (scRNA-seq, microbiome). Eight simulation generators produce correlated biomarkers with controlled properties under exchangeable, AR(1), block, or user-supplied correlation structures with built-in calibration diagnostics. Over a dozen visualization functions are provided. The package is freely available under the MIT license at <https://github.com/vanhungtran/aucmat>.

Availability: **aucmat** v0.2.0, R $\geq$ 4.0.0, MIT license. <https://github.com/vanhungtran/aucmat> (source), <https://vanhungtran.github.io/aucmat> (documentation).

Contact: lucas.tran@ck-care.ch

1 Introduction

Receiver operating characteristic (ROC) analysis is the dominant framework for evaluating continuous biomarkers against binary outcomes in diagnostic medicine, prognostic modeling, and molecular screening [Hanley and McNeil, 1982, DeLong et al., 1988,

Pepe, 2004]. The area under the ROC curve (AUC) summarises the probability that a randomly chosen positive subject has a higher biomarker value than a randomly chosen negative subject.

Modern molecular profiling technologies routinely measure thousands of features on a few hundred subjects. A proteomics panel may contain 5,000 proteins; a transcriptomics experiment 20,000 genes. The core task – identifying which biomarkers discriminate cases from controls – requires computing and testing an AUC for each feature, adjusting for multiplicity, comparing top candidates, and assessing ranking stability. Yet no single widely-used R package integrates these steps.

The R ecosystem provides excellent point solutions. **pROC** [Robin et al., 2011] is the gold standard: it computes DeLong confidence intervals using the Sun & Xu (2014) $\mathcal{O}(N \log N)$ algorithm [Sun and Xu, 2014], tests paired or unpaired AUC differences via `roc.test()`, provides partial AUC with bootstrap inference, binormal and density-smoothed ROC curves, and multiclass AUC via Hand & Till [Hand and Till, 2001]. However, **pROC** operates one biomarker at a time. **ROCR** [Sing et al., 2005] offers 28 performance measures and visualisation but no statistical inference. **precRec** [Saito and Rehmsmeier, 2017] adds fast C++-based precision-recall curves and confidence bands but provides no DeLong inference or multiplicity adjustment. **colAUC** [Tuszynski and Dietze, 2024] computes column-wise AUCs efficiently but provides no inference whatsoever. **dtComb** [Yerlitas et al., 2025] combines two biomarkers via 142 methods but does not perform matrix-wide screening. **cancerclass** [Budczies et al., 2024] builds classifiers from high-dimensional data but lacks DeLong-based screening.

aucmat fills this gap with a *matrix-first* design. The primary function `aucmat(X, y)` accepts $\mathbf{X} \in \mathbb{R}^{n \times p}$ and a binary outcome $\mathbf{y} \in \{0, 1\}^n$, returning point estimates, standard errors, confidence intervals, raw and multiplicity-adjusted p -values, and diagnostic meta-data for every biomarker in one call. Downstream functions for comparison, stability, cross-validation,

panel construction, and power analysis consume the same S3 object, preserving outcome encoding throughout.

2 Methods

2.1 Direction-preserving AUC

For biomarker X and binary outcome $Y \in \{0, 1\}$, the AUC equals the Mann-Whitney U probability:

$$\text{AUC} = P(X_i > X_j \mid Y_i = 1, Y_j = 0) \quad (1)$$

Most software reports $\max(\text{AUC}, 1 - \text{AUC})$, silently reversing the biomarker. `aucmat` preserves the observed direction, reporting AUC_{raw} (may be < 0.5), $\text{AUC}_{\text{strength}} = 0.5 + |\text{AUC}_{\text{raw}} - 0.5|$ (for ranking), and `effect.direction`. A tumour suppressor with $\text{AUC}_{\text{raw}} = 0.15$ discriminates as strongly as an oncogene with 0.85.

2.2 DeLong variance and joint covariance

The DeLong estimator [DeLong et al., 1988, Sun and Xu, 2014] uses placement values:

$$V_{10}(X_i) = \frac{1}{n_0} \sum_{j: Y_j=0} [\mathbf{1}(X_j < X_i) + \frac{1}{2}\mathbf{1}(X_j = X_i)], \quad i \in \mathcal{I}_1 \quad (2)$$

$$V_{01}(X_j) = \frac{1}{n_1} \sum_{i: Y_i=1} [\mathbf{1}(X_i > X_j) + \frac{1}{2}\mathbf{1}(X_i = X_j)], \quad j \in \mathcal{I}_0 \quad (3)$$

Let $S_{10} = (n_1 - 1)^{-1} \sum_{i \in \mathcal{I}_1} (V_{10}(X_i) - \bar{V}_{10})^2$ and $S_{01} = (n_0 - 1)^{-1} \sum_{j \in \mathcal{I}_0} (V_{01}(X_j) - \bar{V}_{01})^2$. The variance estimator is:

$$\widehat{\text{Var}}(\widehat{\text{AUC}}) = \frac{S_{10}}{n_1} + \frac{S_{01}}{n_0} \quad (4)$$

For p biomarkers on common subjects, the $p \times p$ joint covariance is:

$$\widehat{\Sigma} = \frac{1}{n_1} \widehat{\text{Cov}}(\mathbf{V}_{10}) + \frac{1}{n_0} \widehat{\text{Cov}}(\mathbf{V}_{01}) \quad (5)$$

This supports an omnibus Wald test of $H_0 : \text{AUC}_1 = \dots = \text{AUC}_p$. With contrast matrix $\mathbf{C}$ (default: comparing biomarkers $2, \dots, p$ against biomarker 1):

$$W = (\mathbf{C}\hat{\theta})^\top (\mathbf{C}\widehat{\Sigma}\mathbf{C}^\top)^{-1} (\mathbf{C}\hat{\theta}) \sim \chi_r^2 \quad (6)$$

using the effective rank r via eigenvalue tolerance. At $p = 2$, W equals the squared paired-DeLong z -statistic within numerical tolerance.

2.3 Multiplicity and stability

Screening p biomarkers produces p tests. `aucmat` applies Benjamini-Hochberg [Benjamini and Hochberg, 1995], Holm, or Bonferroni correction, reporting adjusted p -values as q -values. For TOST equivalence [Schuirmann, 1987], $p_{\text{equiv}} = \max(p_{\text{lower}}, p_{\text{upper}})$ with adjustment on this combined value.

Bootstrap rank-stability (`auc.stability()`) resamples data B times (stratified by outcome), recomputes the full screening, and reports median/IQR rank distributions, top- k selection probabilities, and pairwise co-selection frequencies [Efron and Tibshirani, 1994].

2.4 Hurdle-AUC for zero-inflated data

Standard AUC assumes continuous biomarkers. On zero-inflated data (scRNA-seq, microbiome), 60%+ of values are zero, and standard AUC degrades to near 0.5. The Hurdle-AUC models zeros and expression as separate processes:

$$\text{Stage 1 (zero hurdle): } P(X_j = 0 \mid Y) = \text{logistic}(\beta_0 + \beta_1 Y) \quad (7)$$

$$\text{Stage 2 (magnitude): } \text{AUC}(X_j \mid X_j > 0) \text{ computed on expressed values} \quad (8)$$

The composite score for observation i is:

$$S_i = \begin{cases} \hat{p}_i & \text{if } X_{ij} = 0 \\ 0.5 \hat{p}_i + 0.5 \tilde{x}_{ij} & \text{if } X_{ij} > 0 \end{cases} \quad (9)$$

where $\hat{p}_i$ is the predicted non-zero probability from Stage 1 and $\tilde{x}_{ij}$ is the normalised biomarker value.

2.5 Multivariate simulation engine

The simulation engine is a core contribution of `aucmat`. We describe it in full mathematical detail, as it underpins the package's reproducibility guarantees and distinguishes it from ad-hoc simulation approaches.

2.5.1 The equal-variance binormal model

Under the standard binormal model [Hanley and McNeil, 1982], for biomarker $k \in \{1, \dots, p\}$:

$$X_k \mid Y = 0 \sim \mathcal{N}(\mu_{k0}, \sigma^2), \quad X_k \mid Y = 1 \sim \mathcal{N}(\mu_{k1}, \sigma^2) \quad (10)$$

The standardised mean separation $\delta_k = (\mu_{k1} - \mu_{k0})/\sigma$ determines the AUC via:

$$\text{AUC}_k = \Phi\left(\frac{\delta_k}{\sqrt{2}}\right) \quad (11)$$

where $\Phi(\cdot)$ is the standard normal CDF. Inverting gives the closed-form calibration used by all `aucmat` simulators:

$$\delta_k = \sqrt{2} \cdot \Phi^{-1}(\text{AUC}_k) \quad (12)$$

For $\text{AUC} = 0.8$, $\delta \approx 1.19$; for $\text{AUC} = 0.9$, $\delta \approx 1.81$. The class means are centred so the n -weighted grand mean is zero:

$$\mu_{k0} = \delta_k \cdot \left(-\frac{n_1}{n}\right), \quad \mu_{k1} = \delta_k \cdot \left(\frac{n_0}{n}\right) \quad (13)$$

2.5.2 Class-conditional multivariate normal (simulate_auc_matrix)

For p biomarkers with correlation matrix $\mathbf{R}$, the class-conditional distribution generalises Eq. (10) to the multivariate case:

$$\mathbf{X} \mid Y = 0 \sim \mathcal{N}_p(\boldsymbol{\mu}_0, \mathbf{R}), \quad \mathbf{X} \mid Y = 1 \sim \mathcal{N}_p(\boldsymbol{\mu}_1, \mathbf{R}) \quad (14)$$

where $\boldsymbol{\mu}_0 = (\mu_{10}, \dots, \mu_{p0})^\top$ and $\boldsymbol{\mu}_1 = (\mu_{11}, \dots, \mu_{p1})^\top$ follow Eq. (13). Both classes share the same $\mathbf{R}$.

The outcome Y is fixed *before* any biomarker is drawn. In `outcome_mode = "exact"`, $n_1 = \lfloor n\pi \rfloor$ observations are assigned $Y = 1$ with randomised order; in `"bernoulli"`, each $Y_i \sim \text{Bernoulli}(\pi)$. When Y is user-supplied, its exact values and row order are preserved.

The matrix $\mathbf{R}$ is built from a named structure:

- **exchangeable**: $R_{ij} = \rho$ for all $i \neq j$. Requires $\rho > -1/(p-1)$ for positive definiteness.
- **AR(1)**: $R_{ij} = \rho^{|i-j|}$.
- **block**: $R_{ij} = \rho_w$ if i, j in same block, ρ_b otherwise. Block sizes must sum to p .
- **user**: R_{ij} directly from the supplied matrix.

When $\mathbf{R}$ is not positive definite, `feasibility = "nearest"` projects to the closest valid correlation matrix via eigenvalue clipping followed by re-normalisation to unit diagonal. The adjustment is reported in `$feasibility` with Frobenius norm and maximum element-wise change.

After generation, each biomarker is standardised to zero mean and unit variance (an affine transform that preserves both AUC and Pearson correlation). Achieved correlation is computed on within-class centred residuals to avoid mixture-induced inflation (Simpson's paradox).

2.5.3 Gaussian copula with perturbation (simulate_auc_copula)

This simulator separates correlation control from discriminative signal:

Phase 1 – Gaussian Copula. For each class independently, n_c draws are taken from $\mathcal{N}_{p+1}(\mathbf{0}, \mathbf{R}_{\text{work}})$,

where $\mathbf{R}_{\text{work}}$ is a $(p+1) \times (p+1)$ matrix combining the biomarker correlation $\mathbf{R}$ (upper-left $p \times p$) with latent outcome correlations $\boldsymbol{\rho}$ (last row/column). The latent correlations are approximated from target AUCs:

$$\rho_k^{\text{latent}} = \frac{\delta_k}{\sqrt{1 + \delta_k^2}} \quad (15)$$

Each column is rank-transformed to uniform quantiles, then probit-transformed to standard normal scores. This copula step preserves rank correlations exactly regardless of marginal distributions.

Phase 2 – Iterative Perturbation. After Phase 1, each biomarker's empirical AUC, $\widehat{\text{AUC}}_k$, is compared to its target. The gap $\Delta_k = \text{AUC}_k - \widehat{\text{AUC}}_k$ is translated to a mean shift via the binormal calibration:

$$\Delta\delta_k = \sqrt{2} \left[\Phi^{-1}(\text{AUC}_k) - \Phi^{-1}(\widehat{\text{AUC}}_k) \right] \quad (16)$$

The shift $\Delta\delta_k \cdot \sigma_k \cdot \eta^t$ is applied class-conditionally (positive class receives $+(n_0/n)$ times the shift, negative class $-(n_1/n)$), where $\eta \in (0, 1)$ is the step decay and t the iteration index. Perturbation stops when $|\Delta_k| < \tau$ for all k or after T_{max} iterations.

2.5.4 Hurdle simulator (simulate_hurdle_auc)

Generates zero-inflated biomarkers via a two-stage process. For biomarker j :

Stage 1 – Zero hurdle. Logistic parameters are solved from user-specified per-class zero rates:

$$\beta_0 = \text{logit}(p_0), \quad p_0 = P(X_j = 0 \mid Y = 0) \quad (17)$$

$$\beta_1 = \text{logit}(p_1) - \text{logit}(p_0), \quad p_1 = P(X_j = 0 \mid Y = 1) \quad (18)$$

For each subject, $Z_{ij} \sim \text{Bernoulli}(\text{logistic}(\beta_0 + \beta_1 Y_i))$ determines whether $X_{ij} = 0$.

Stage 2 – Magnitude. For observations with $Z_{ij} = 0$ (expressed), values are drawn from a log-normal distribution with class-conditional means calibrated to achieve the target non-zero AUC via Eq. (12):

$$X_{ij} \mid X_{ij} > 0, Y_i = c \sim \text{LogNormal}(\mu_c, \sigma^2) \quad (19)$$

where $\mu_1 - \mu_0$ is set by the binormal shift. Values are clipped to $[0, \infty)$.

2.5.5 Other simulators

`generate_data_probit()` uses a latent probit formulation where $\mathbf{X}$ and a latent $Y^* \sim \mathcal{N}(0, 1)$ are jointly $\mathcal{N}_{p+1}(\mathbf{0}, \boldsymbol{\Sigma})$, with $Y = \mathbf{1}\{Y^* > \Phi^{-1}(1 - \pi)\}$. The full $(p+1) \times (p+1)$ correlation matrix simultaneously

controls both biomarker-biomarker and biomarker-outcome associations. The latent correlations $\rho_k = \text{Cor}(X_k, Y^*)$ are calibrated numerically via bisection over a large calibration sample ($n_{\text{cal}} = 100,000$).

`generate_data_analytical()` constructs biomarkers sequentially from highest to lowest AUC. For $k \geq 2$, X_k is built as a linear combination of $X_1, \dots, X_{k-1}$ (weighted by $\mathbf{R}_{k,1:k-1} \mathbf{R}_{1:k-1,1:k-1}^{-1}$ to achieve target correlations) plus an independent residual carrying the remaining mean separation. This is fast (closed-form, no calibration loop) but links AUC and correlation.

`generate_auc_vector()` constructs a single score with exact empirical AUC via rank assignment. Given n_1 positives and n_0 negatives, achievable AUCs form a discrete grid with step $\Delta_{\text{AUC}} = 1/(n_1 n_0)$. The required rank sum is solved and normal quantile scores are assigned.

`generate_auc_cor_vector()` extends this with a strictly increasing Box-Cox transform tuned to achieve a target Pearson correlation without changing the AUC (since monotone transforms preserve ranks).

`simulate_auc_correlation()` repeatedly draws from the binormal model and records achieved (AUC, r) pairs to characterise their sampling distribution.

2.5.6 Simulation validation

`validate_simulation()` repeats a specification across R independent replicates with deterministic seed streams (seed + $r - 1$). For each replicate it records achieved AUCs and correlations, then reports:

$$\text{Bias}_k = \frac{1}{R} \sum_{r=1}^R (\widehat{\text{AUC}}_{kr} - \text{AUC}_k) \quad (20)$$

$$\text{RMSE}_k = \sqrt{\frac{1}{R} \sum_{r=1}^R (\widehat{\text{AUC}}_{kr} - \text{AUC}_k)^2} \quad (21)$$

$$\text{MCSE}_k = \frac{\text{SD}_k}{\sqrt{R}} \quad (22)$$

plus target-interval hit rates for both AUCs and correlations. A bias large relative to MCSE signals a systematic calibration error.

3 Implementation

`aucmat` is implemented in pure R (6,200 lines, 34 source files). The primary function performs six steps: (1) input validation, (2) outcome normalisation (accepting logical, 0/1, factor, or character), (3) matrix-wide AUC via Mann-Whitney rank formulation, (4) DeLong or bootstrap inference, (5) multiplicity adjustment, (6) assembly with ranking and warnings.

All functions return S3 objects with `print()`, `summary()`, `plot()`, and `as.data.frame()` methods. Structured error classes allow programmatic handling. An RNG-safety mechanism (`with_seed()`) restores `.Random.seed` on exit.

Dependencies: pROC (≥ 1.18), mvtnorm, ggplot2, rlang. Optional Suggests: GGally, glmnet, matrixStats, knitr, rmarkdown, ggrepel.

4 Usage

Basic screening:

```
fit <- aucmat(X, y, ci = "delong", adjust = "BH")
print(fit); summary(fit)
plot_auc_rank(fit, n_label = 4)
```

Paired comparisons with hypothesis testing:

```
compare_auc(fit, X, y, reference = "X1") # superiority
compare_auc(fit, X, y, reference = "X1",
  hypothesis = "noninferiority", margin = 0.05)
compare_auc(fit, X, y, reference = "X1",
  hypothesis = "equivalence", margin = 0.15)
compare_auc_global(fit, X, y) # H0: all AUCs equal
```

Hurdle-AUC for zero-inflated data:

```
sim <- simulate_hurdle_auc(n = 400, prevalence = 0.3,
  target_hurdle_auc = c(0.85, 0.72),
  zero_rate_neg = c(0.55, 0.80),
  zero_rate_pos = c(0.25, 0.70))
fit <- hurdle_auc(as.matrix(sim$data[,1:2]), sim$data$strut)
# Standard AUC ~0.53 -> Hurdle AUC ~0.90
```

Cross-validation, panel construction, power analysis:

```
cv <- cv_aucmat(X, y, n_folds = 5, seed = 42)
panel <- fit_auc_panel(X[,top4], y, method = "ridge")
power_auc_matrix(target_auc = 0.70, power = 0.80,
  prevalence = 0.3)
```

5 Benchmarking

Table 1: Empirical coverage of DeLong 95% confidence intervals, averaged across 27 simulation conditions (200 replicates each).

AUC	$n = 100$	$n = 300$	$n = 500$
0.65	0.898	0.937	0.957
0.75	0.915	0.943	0.933
0.85	0.883	0.923	0.905

Screening throughput. On a desktop workstation (AMD Ryzen 3.5 GHz, 32 GB RAM, R 4.5.1, Windows 11), `aucmat` screens 100, 1,000, and 10,000 biomarkers ($n = 500$, DeLong inference) in 0.54, 5.02, and 59.76 seconds. The equivalent per-column `pROC` loop requires 0.68, 6.42, and an estimated 64 seconds.

Simulation calibration. Across 27 conditions ($n = 100, 300, 500$; prevalence = 0.2, 0.3, 0.5; AUC = 0.65, 0.75, 0.85; 200 reps each), DeLong coverage ranged from 0.898 to 0.957, converging to 0.95 as n increased (Table 1). Power exceeded 0.85 for AUC ≥ 0.65 at $n = 100$ and reached 1.0 for AUC ≥ 0.75 at $n \geq 300$.

6 Comparison with existing packages

Five capabilities are unique to `aucmat`: (i) matrix-first screening with direction-preserving AUC and multiplicity adjustment; (ii) non-inferiority and equivalence tests for paired AUCs; (iii) omnibus Wald test for three or more correlated AUCs; (iv) bootstrap rank-stability analysis; and (v) Hurdle-AUC for zero-inflated data.

For a single biomarker with smoothed ROC curves, `pROC` remains the standard. For speed on $> 10^5$ columns, `colAUC` is recommended. For precision-recall of imbalanced data, `precRec` is preferred. For combining two biomarkers, `dtComb` provides the most exhaustive methods. For end-to-end biomarker screening with full inference, `aucmat` is the recommended tool.

7 Discussion

`aucmat` provides a unified, statistically principled framework for biomarker screening. Its matrix-first design reduces the code burden from dozens of scripting steps to a few function calls.

Limitations. The current version (0.2.0) is restricted to binary outcomes for the primary workflow (multiclass is experimental). DeLong requires common subjects; unpaired comparisons are not implemented (`pROC` provides this). Partial AUC is an internal engine not yet exposed in `aucmat()`. The DeLong estimator is slightly liberal at small n (coverage ≈ 0.90 at $n = 100$), consistent with its finite-sample behaviour; bootstrap intervals are recommended when $n < 100$ per class.

Future work. v0.3: nested cross-validation for penalised panel models. v0.4: multiclass AUC with one-versus-rest, one-versus-one, and Hand-Till estimators. v0.5: cumulative/dynamic time-dependent AUCs for survival outcomes [Heagerty et al., 2000,

Blanche et al., 2013]. v0.6: sparse-matrix and on-disk backends for genomics-scale data.

Acknowledgements

We thank The Kühne Foundation for financial support. Xavier Robin and colleagues are acknowledged for `pROC`, which provided the statistical foundation and ROC curve objects. Hadley Wickham’s `ggplot2` and the `mvtnorm` authors (Genz, Bretz, et al.) are gratefully acknowledged. The Sun & Xu (2014) fast DeLong algorithm implemented in `pROC` since v1.9 informs `aucmat`’s variance computation.

Funding

This work was supported by The Kühne Foundation.

References

- Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300, 1995.
- Paul Blanche, Jean-François Dartigues, and Hélène Jacqmin-Gadda. Estimating and comparing time-dependent areas under receiver operating characteristic curves for censored event times with competing risks. *Statistics in Medicine*, 32(30):5381–5397, 2013.
- Jan Budczies, Daniel Kosztyla, Christian von Törne, Albrecht Stenzinger, Silvia Darb-Esfahani, Manfred Dietel, Carsten Denkert, and Frederick Klauschen. `cancerclass`: Development and validation of diagnostic tests from high-dimensional molecular data, 2024. R package version 1.49.0, Bioconductor.
- Elizabeth R DeLong, David M DeLong, and Daniel L Clarke-Pearson. Comparing the areas under two or more correlated receiver operating characteristic curves: a nonparametric approach. *Biometrics*, 44(3):837–845, 1988.
- Bradley Efron and Robert J Tibshirani. *An Introduction to the Bootstrap*. CRC Press, Boca Raton, FL, 1994.
- David J Hand and Robert J Till. A simple generalisation of the area under the ROC curve for multiple class classification problems. *Machine Learning*, 45(2):171–186, 2001.

Table 2: Feature comparison of R packages for ROC analysis and biomarker screening. + = supported; -- = not supported. Verified against CRAN documentation as of July 2026.

Feature	aucmat	pROC	precrc	ROCR	colAUC	dtComb
Matrix-first screening	+	—	—	—	+	—
Direction-preserving AUC	+	—	—	—	—	—
DeLong CI	+	+	—	—	—	—
Stratified bootstrap CI	+	+	—	—	—	+
Multiplicity adjustment	+	—	—	—	—	—
One-sided tests	+	+	—	—	—	—
Paired comparisons	+	+	—	—	—	+
Non-inferiority/equivalence	+	—	—	—	—	—
Omnibus Wald (3+ AUCs)	+	—	—	—	—	—
Bootstrap rank stability	+	—	—	—	—	—
Partial AUC	+	+	+	—	—	—
Multivariate simulation	+	—	—	—	—	—
Simulation diagnostics	+	—	—	—	—	—
Cross-validated AUC	+	—	—	—	—	+
Panel (ridge/lasso)	+	—	—	—	—	+
Hurdle-AUC (zero-inflated)	+	—	—	—	—	—
Precision-Recall curves	+	—	+	—	—	—
Power / sample size	+	+	—	—	—	—
Multiclass (Hand-Till)	+	+	—	—	+	—
ROC smoothing	+	+	—	—	—	—
Group-stratified screening	+	—	—	—	—	—

- James A Hanley and Barbara J McNeil. The meaning and use of the area under a receiver operating characteristic (ROC) curve. *Radiology*, 143(1):29–36, 1982.
- Patrick J Heagerty, Thomas Lumley, and Margaret S Pepe. Time-dependent ROC curves for censored survival data and a diagnostic marker. *Biometrics*, 56(2):337–344, 2000.
- Margaret S Pepe. *The Statistical Evaluation of Medical Tests for Classification and Prediction*. Oxford University Press, New York, 2004.
- Xavier Robin, Natacha Turck, Alexandre Hainard, Natalia Tiberti, Frédérique Lisacek, Jean-Charles Sanchez, and Markus Müller. pROC: an open-source package for R and S+ to analyze and compare ROC curves. *BMC Bioinformatics*, 12:77, 2011. doi: 10.1186/1471-2105-12-77.
- Takaya Saito and Marc Rehmsmeier. precrc: fast and accurate precision-recall and ROC curve calculations in R. *Bioinformatics*, 33(1):145–147, 2017. doi: 10.1093/bioinformatics/btw570.
- Donald J Schuirmann. A comparison of the two one-sided tests procedure and the power approach for assessing the equivalence of average bioavailability. *Journal of Pharmacokinetics and Biopharmaceutics*, 15(6):657–680, 1987.
- Tobias Sing, Oliver Sander, Niko Beerenwinkel, and Thomas Lengauer. ROCR: visualizing classifier performance in R. *Bioinformatics*, 21(20):3940–3941, 2005. doi: 10.1093/bioinformatics/bti623.
- Xu Sun and Weichao Xu. Fast implementation of DeLong’s algorithm for comparing the areas under correlated receiver operating characteristic curves. *IEEE Signal Processing Letters*, 21(11):1389–1393, 2014. doi: 10.1109/LSP.2014.2337313.
- Jarek Tuszynski and Michael Dietze. caTools: Tools: moving window statistics, GIF, Base64, ROC AUC, etc., 2024. R package version 1.18.3.
- Serra Ilayda Yerlitas, Serra Bersan Gengec, Necla Kochan, Gozde Erturk Zararsiz, Selcuk Korkmaz, and Gokmen Zararsiz. dtComb: A comprehensive R library and web tool for combining diagnostic tests. *The R Journal*, 17(1):1–20, 2025. doi: 10.32614/RJ-2025-036.